

Trust Rust! (For Safe Systems Programming)

Pramode C.E

March 2, 2017

Systems programming concepts for application developers

The Stack

```
// a0.c
```

```
int main()
{
    int a = 1, b = 2;
    return 0;
}
```

The Stack




Figure 1: The C stack

Buffer overflow

```
// a1.c
int main()
{
    int i=1;
    int a[4];
    int j=2;

    a[4] = 5; // bug
    a[5] = 6; // bug
    a[10000] = 7; // bug
}
```

Buffer overflow




Figure 2: Buffer overflow diagram

Pointers in C

```
// a2.c  
int main()  
{  
    int a = 10;  
    int *b;  
  
    b = &a;  
    *b = 20;  
}
```

Pointers in C




Figure 3: How pointer variables are stored in memory

Stack Frames - deallocations and allocations

```
// a3.c
void fun2() { int e=5, f=6; }
void fun1() { int c=3, d=4; }

int main()
{
    int a=1, b=2;
    fun1(); fun2();
    return 0;
}
```

Stack Frames - allocations and deallocations




Figure 4: Multiple stack frames

Stack Frames - allocations and deallocations




Figure 5: Multiple stack frames

Dangling Pointers

```
// a4.c
void fun2() { int m = 1; int n = 2; }
int* fun1() {
    int *p; int q = 0;
    p = &q; return p; // bug
}

int main() {
    int *a, b; a = fun1();
    *a = 10; fun2();
    b = *a;
}
```

Dangling pointers




Figure 6: Dangling pointer

Null pointer dereferencing

```
// a6.c
#include <strings.h>
#include <stdio.h>
int main()
{
    char s[100] = "hello";
    char *p;
    p = index(s, 'f');
    *p = 'a'; // bug!
    return 0;
}
```

Heap allocation - malloc and free

```
// a7.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
    free(c);
}
int main()
{
    fun();
}
```

Heap allocation - malloc and free




Figure 7: A stack location pointing to a heap location

Memory leaks

```
// a8.c
#include <stdlib.h>
void fun()
{
    char *c;
    c = malloc(10*sizeof(char));
    /* do some stuff here */
}
int main()
{
    fun(); // bug! memory leak.
}
```

Use-after-free

```
// a9.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    c[0] = 'A'; //bug! user-after-free
}
```

Double free

```
// a10.c
#include <stdlib.h>
void fun(char *t) {
    /* do some stuff here */
    free(t);
}
int main() {
    char *c;
    c = malloc(10 * sizeof(char));
    fun(c);
    free(c); //bug! double free
}
```

John Regehr, on C

When tools like the bounds checking GCC, Purify, Valgrind, etc. first showed up, it was interesting to run a random UNIX utility under them. The output of the checker showed that these utility programs, despite working perfectly well, executed a ton of memory safety errors such as use of uninitialized data, accesses beyond the ends of arrays, etc. Just running grep or whatever would cause tens or hundreds of these errors to happen.

From: <http://blog.regehr.org/archives/226>

John Regehr, on C

More and more, I'm starting to wonder how safety-critical code can continue being written in C.

A comment on: <http://blog.regehr.org/archives/232>

Rust and Memory Safety

- ▶ Let's get into the real core of Rust!

Scope

```
// a12-3.rs  
fn main() {  
    let x = 10;  
    {  
        let y = 20;  
    }  
    println!("x={}, y={}", x, y);  
}
```

Ownership

```
// a13.rs  
fn main() {  
    let v = vec![10, 20, 30];  
    println!("{:?}", v);  
}  
// how is v deallocated?
```


Ownership




Figure 8: Memory representation of a vector

Ownership

```
// a14.rs  
fn fun1() {  
    let v = vec![10 ,20 ,30];  
} // how is v deallocated?  
fn main() {  
    fun1();  
}
```

Ownership

```
// a15.rs  
fn main() {  
    let v1 = vec![10 ,20 ,30];  
    let v2 = v1;  
    println!("{:?}", v2);  
}  
// do we have a double free here?
```

Ownership

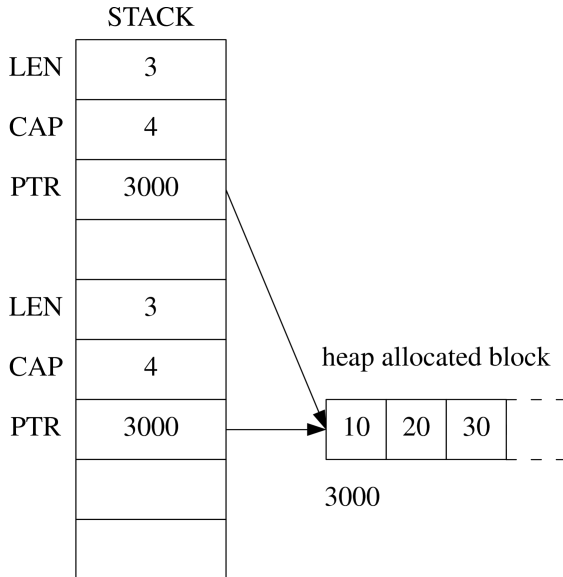


Figure 9: Two pointers

Ownership

```
// a15-1.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
}  
// do we have a double free here?
```

Ownership

```
// a16.rs  
fn main() {  
    let v1 = vec![10, 20, 30];  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v2);  
}  
// what happens if we try to acces the  
// vector through v1?
```

Move semantics

```
// a17.rs  
fn main() {  
    let v1 = vec![1,2,3];  
  
    let mut v2 = v1;  
    v2.truncate(2);  
    println!("{:?}", v1);  
}
```

Move semantics

```
// a15-2.rs  
fn fun(v2: Vec<i32>) {  
    println!("{:?}", v2);  
}  
fn main() {  
    let v1 = vec![10, 20, 30];  
    fun(v1);  
    println!("{:?}", v1);  
}
```


Move semantics

```
// a18.rs  
fn main() {  
    let a = (1, 2.3);  
    let b = a;  
    println!("{}", a);  
}
```

Move semantics

```
// a19.rs  
fn main() {  
    let a = (1, 2.3, vec![10,20]);  
    let b = a;  
    println!("{:?}", a);  
}
```

Memory safety without garbage collection

- ▶ Languages like Python, Java etc achieve memory safety at run time through garbage collection.
- ▶ Rust achieves memory safety at compile time by static type analysis.
- ▶ Ownership + move semantics has some interesting properties which makes them suitable for general resource management (not just memory).

Garbage Collection

```
# a20.py  
a = [10, 20, 30]  
a.append(40)  
print a
```

Garbage Collection

```
# a21.py  
a = [10, 20, 30]  
b = a  
b.append(40)  
print a # what does this print?
```

Garbage Collection




Figure 10: References in Python

Garbage Collection




Figure 11: Reference Counting

Garbage Collection

```
# a22.py
```

```
a = [10, 20, 30]
```

```
b = a # refcount is 2
```

```
a = "hello" # refcount is 1
```

```
b = "world" # refcount drops to zero, deallocate
```


Ownership / Move: Limitations

```
// a26.rs
fn vector_sum(v: Vec<i32>) -> i32 {
    //assume v is always a 3 elemnt vector
    v[0] + v[1] + v[2]
}
fn main() {
    let v = vec![1,2,3];
    let s = vector_sum(v);
    println!("{}",s);
}
```

Ownership / Move: Limitations

```
// a27.rs  
fn vector_sum(v: Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
  
fn vector_product(v: Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(v);  
    let p = vector_product(v);  
    println!("{}",p);  
}  
  
// does this code compile?
```

Immutable Borrow

```
// a28.rs  
fn vector_sum(v: &Vec<i32>) -> i32 {  
    v[0] + v[1] + v[2]  
}  
fn vector_product(v: &Vec<i32>) -> i32 {  
    v[0] * v[1] * v[2]  
}  
fn main() {  
    let v = vec![1,2,3];  
    let s = vector_sum(&v);  
    let p = vector_product(&v);  
    println!("v={:?}, s={}, p={}", v, s, p);  
}
```

Immutable Borrow

```
// a29.rs  
fn main() {  
    let v = vec![1,2,3];  
    let t1 = &v;  
    let t2 = &v;  
    println!("{}", t1[0], t2[0], v[0]);  
}  
// any number of immutable borrows are ok!
```

Immutable Borrow

```
// a30.rs  
fn change(t1: &Vec<i32>) {  
    t1[0] = 10;  
}  
fn main() {  
    let mut v = vec![1,2,3];  
    change(&v);  
}  
// Does the program compile?
```

Mutable Borrow

```
// a31.rs  
fn change(t1: &mut Vec<i32>) {  
    t1[0] = 10;  
}  
fn main() {  
    let mut v = vec![1,2,3];  
    change(&mut v);  
    println!("{:?}", v);  
}
```

A use-after-free bug

```
// a32.c
#include <stdlib.h>
int main()
{
    char *p = malloc(10 * sizeof(char));
    char *q;

    q = p + 2;
    free(p);
    *q = 'A'; // bug!
    return 0;
}
```

Vector allocation in Rust

```
// a33.rs  
fn main() {  
    let mut a = vec![];  
    a.push(1); a.push(2);  
    a.push(3); a.push(4);  
    a.push(5);  
  
    println!("{:?}", a);  
}
```


Vector allocation in Rust/C++




Figure 12: Growing a vector

A use-after-free bug in C++

```
// a33-1.cpp
#include <vector>
#include <iostream>
using namespace std;
int main()
{
    vector<int> v;
    int *p;
    v.push_back(1);
    p = &v[0];
    v.push_back(2);
    *p = 100; // bug!
    cout << v[0] << endl;
}
```

A use-after-free bug in Rust?

```
// a34.rs  
fn main() {  
    let mut v = vec![10, 20, 30, 40];  
    let p1 = &v[1];  
    v.push(50);  
    // bug if we try to use p1  
    // does this code compile?  
}
```

Borrowing Rules

- ▶ Any number of immutable borrows can co-exist.
- ▶ A mutable borrow can not co-exist with other mutable or immutable borrows.
- ▶ The “borrow checker” checks violations of these rules at compile time.

Borrow checker limitations

- ▶ The borrow checker gives you safety by rejecting ALL unsafe programs.
- ▶ But it is not perfect in the sense it rejects safe programs also; “fighting the borrow checker” is a common sporting activity among Rust programmers :)
- ▶ There are plans to improve the situation:
<http://smallcultfollowing.com/babysteps/blog/2017/03/01/nested-method-calls-via-two-phase-borrowing/>

Borrow checker limitations - an example

```
// a35.rs  
fn main() {  
    let mut v = vec![10,20,30];  
    v.push(v.len());  
}  
// this will not compile
```

Borrow checker limitations - an example

```
// a36.rs  
// Same as a35.rs  
fn main() {  
    let mut v = vec![10,20,30];  
    let tmp0 = &v;  
    let tmp1 = &mut v;  
    let tmp2 = Vec::len(tmp0); //v.len()  
  
    Vec::push(tmp1, tmp2); // v.push(tmp2)  
}
```

Lifetimes

```
// a37.rs
fn main() {
    let ref1: &Vec<i32>;
    {
        let v = vec![1, 2, 3];
        ref1 = &v;
    }
    // v gets deallocated as it goes out of
    // the scope. What about ref1? Do we have
    // a "dangling pointer" here?
}
```


Lifetimes

```
// a38.rs
fn foo() -> Vec<i32> {
    let v = vec![1, 2, 3];
    v // transfer ownership to caller
}

fn main() {
    let p = foo();
    println!("{:?}", p);
}
```

Lifetimes

```
// a39.rs
fn foo() -> &Vec<i32> {
    let v = vec![1, 2, 3];
    &v // Will this compile?
}

fn main() {
    let p = foo();
}
```

Explicit Lifetime Annotations

```
// a40.rs
fn foo(v1: &Vec<i32>, v2: &Vec<i32>) -> &i32 {
    &v1[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
        // How does the compiler know, just by looking at
        // the signature of "foo", that the reference
        // returned by "foo" will live as long as "p"?
    }
}
```

Explicit Lifetime Annotations

```
// a41.rs
fn foo<'a, 'b>(v1: &'a Vec<i32>,
               v2: &'b Vec<i32>) -> &'a i32 {

    &v1[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
    }
}
```

Explicit Lifetime Annotations

```
// a42.rs
fn foo<'a, 'b>(v1: &'a Vec<i32>,
               v2: &'b Vec<i32>) -> &'b i32 {

    &v2[0]
}

fn main() {
    let v1 = vec![1, 2, 3];
    let p:&i32;
    {
        let v2 = vec![4, 5, 6];
        p = foo(&v1, &v2);
    }
}
```

Unsafe

```
// a43.rs
fn main() {
    // a is a "raw" pointer initialized to 0
    let a: *mut u32 = 0 as *mut u32;

    *a = 0;
}
```

Unsafe

```
// a44.rs  
fn main() {  
    let a: *mut u32 = 0 as *mut u32;  
  
    unsafe {  
        *a = 0;  
    }  
}
```

The C FFI

[Demo program in `code/memory-safety/c-ffi`]

Fearless Concurrency!

[Folders: code/fearless-concurrency]

Interesting projects using Rust

- ▶ Servo, from Mozilla. The next-gen browser engine.
- ▶ Redox OS (<https://www.redox-os.org/>), an Operating System being written from scratch in Rust.
- ▶ ripgrep (<https://github.com/BurntSushi/ripgrep>), a fast text search tool.
- ▶ rocket.rs - a powerful web framework.
- ▶ More: <https://github.com/kud1ing/awesome-rust>

Documentation / Learning Resources

- ▶ Official Rust book (<http://rust-lang.github.io/book/>). The second edition is far better, even though it is incomplete.
- ▶ Upcoming O'Reilly book:
<http://shop.oreilly.com/product/0636920040385.do>
- ▶ <http://intorust.com/> (screencasts for learning Rust)
- ▶ <http://exercism.io>
- ▶ Various Conference Videos

Contact Me

- ▶ Email: mail@pramode.net
- ▶ Twitter: [@pramode_ce](https://twitter.com/pramode_ce)

Thank you!

- ▶ Thank You! And, Happy Hacking with Rust!